

X86 Data Arrays and Structures

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Basic data types

■ Integral

- Stored and operated on in general (integer) registers
- Signed vs. unsigned depends on instructions used
- Example:
 - **byte** (size: 1 bytes, appendix: `b`, in C/Java: `char`)
 - **word** (size: 2 bytes, appendix: `w`, in C/Java: `short`)
 - **double word** (size: 4 bytes, appendix: `l`, in C/Java: `int`)
 - **quad word** (size: 8 bytes, appendix: `q`, in C/Java: `long` (x86-64))

■ Floating point

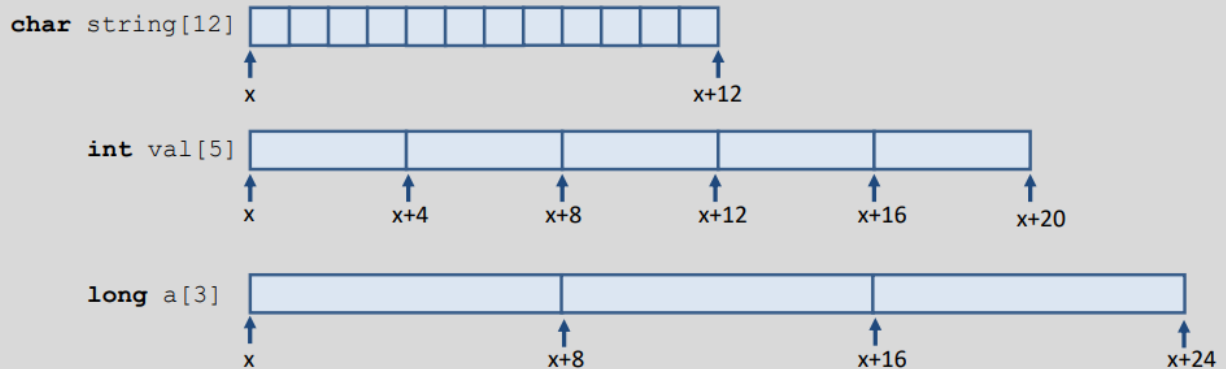
- Stored and operated on floating point registers

Array allocation


Basic Principle

T $A[L];$

- Array of data type T and length L
- Contiguously allocated region of $L * \text{sizeof}(T)$ bytes in memory



Array Accessing Example



```
int get_digit(int val[], int dig)
{
    return val[dig];
}
```

```
get_digit:
    movslq    %rsi, %esi
    movl      (%rdi,%rsi,4), %eax
    ret
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at $4 * \%rsi + \%rdi$
- Use memory reference $(\%rdi, \%rsi, 4)$.

Array Loop Example

```
void digit_inc(int val[], int len){
    int i;
    for (i = 0; i < len; i++)
        val[i]++;
    return;
}
```

```
void digit_inc(int val[], int len){
    if (len>0){
        int *ptr = val;
        int *vend = val + len;
        do {
            *ptr = *ptr + 1;
            ptr++;
        } while (ptr != vend);
    }
    return;
}
```

```
digit_inc:
    testl    %esi, %esi
    jle      .L1
    movq     %rdi, %rax
    leal     -1(%rsi), %edx
    leaq     4(%rdi,%rdx,4), %rdx
.L3:
    addl     $1, (%rax)
    addq     $4, %rax
    cmpq     %rdx, %rax
    jne      .L3
.L1:
    rep ret
```


Nested Arrays

Nested Arrays

```
typedef int zip_dig[5];
zip_dig pgh[4] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```

- “**zip_dig pgh[4]**” equivalent to “**int pgh[4][5]**”
- Variable **pgh**: array of 4 elements, allocated contiguously
- Each element is an array of 5 **int**’s, allocated contiguously
- “**Row-Major**” ordering of all elements guaranteed

Recall, $T\ A[L]$ is an array with L elements of type T .



Multidimensional (nested) Arrays

Declaration

`T A[R][C];`

- 2D array of data type T
- R rows, C columns
- Type T elements require K bytes


$A[0][0]$...	$A[0][C-1]$
$A[1][0]$...	$A[1][C-1]$
...		...
$A[R-1][0]$...	$A[R-1][C-1]$

Array size

- $R * C * K$ bytes

Arrangement

- Row-Major ordering



(Spring'24)

COMP4005: Introduction to Computer Architecture


8

Nested Array Row Access

Row vectors

- $T A[R][C]$: $A[i]$ is array of C elements
- Each element of type T requires K bytes
- Starting address $A + i * (C * K)$

`int A[R][C];`



Nested Array Row Access Mode

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
zip_dig pgh[4] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```

```
leaq    (%rdi,%rdi,4), %rax    # 5 * index
leaq    pgh(,%rax,4), %rax     # pgh + (20 * index)
```

■ Row vector:

- `pgh[index]` is array of 5 int's,
- starting address `pgh+20*index`

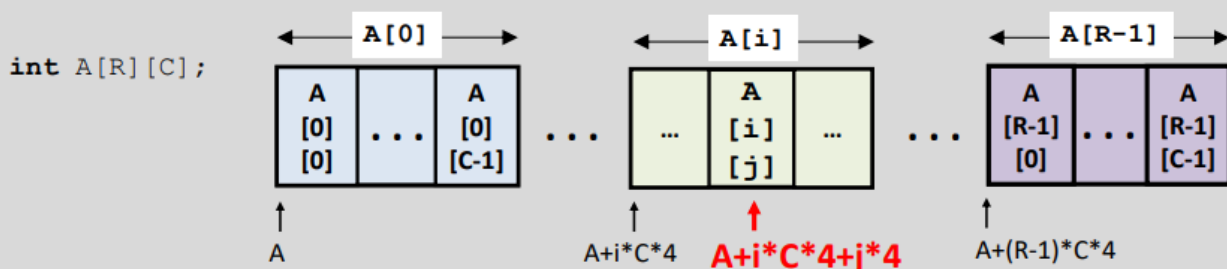
■ x86_64 code:

- Computes and returns address
- Compute as `pgh+4*(index+4*index)`

Nested Array Element Access

■ Array Elements:

- `A[i][j]` is element of type T , which requires K bytes
- Address = $A + i * (C * K) + j * K$
 $= A + (i * C + j) * K$



Nested Array Element Access Code

```
int get_pgh_digit(int index, int digit ){
    return pgh[index][digit];
}
```

```
leaq    (%rdi, %rdi, 4), %rax    # 5 * index
addq    %rax, %rsi              # digit + 5 * index
movl    pgh(, %rsi, 4), %eax    # *(pgh + 4*(digit + 5*index))
```

■ Array Elements:

- `pgh[index][digit]` is `int` and `sizeof(int)=4`
- Address: `pgh + 5*4*index + 4*digit`

■ Assembly Code:

- Computes address as: `pgh + ((index+4*index) + digit)*4`
- `movl` performs memory reference

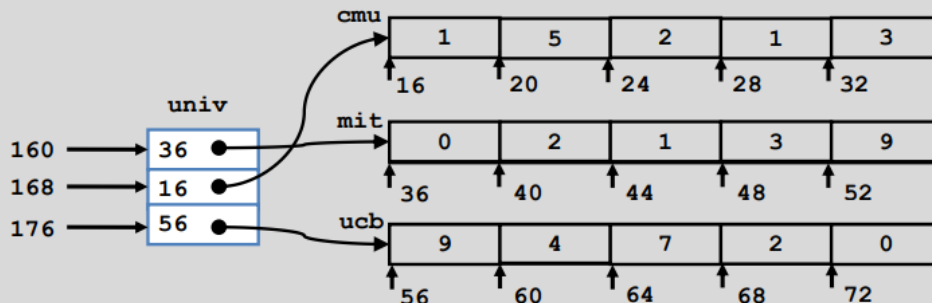
Multi-Level Array Example

Multi-Level Array Declaration(s):

```
int cmu[5] = {1, 5, 2, 1, 3};
int mit[5] = {0, 2, 1, 3, 9};
int ucb[5] = {9, 4, 7, 2, 0};
```

```
int* univ[3] = {mit, cmu, ucb};
```


- Variable `univ` denotes array of 3 elements
- Each element is a pointer: 8 bytes
- Each pointer points to array of `ints`



Note: this is how Java represents multi-dimensional arrays!

Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig ){
    return univ[index][dig];
}
```



```
salq    $2, %rsi          # rsi = 4*dig
addq    univ(,%rdi,8), %rsi # rsi = univ + 8 *index + 4*dig
movl    (%rsi), %eax       # return univ[index][dig] element value
```


Computation (x86_64)

- Element access `Mem[Mem[univ+8*index]+4*dig]`
- Must do **two memory reads**:
 - First get pointer to row array
 - Then access element within array
- But, allows inner arrays to be **different lengths** (although not in this example)

Array Element Accesses

Nested Array


```
int get_pgh_digit
(int index, int digit ){
    return pgh[index][digit];
}
```



`Mem[pgh+20*index+4*digit]`

Multi-level Array

```
int get_univ_digit
(int index, int digit ){
    return univ[index][digit];
}
```



`Mem[Mem[univ+8*index]+4*dig]`

Access looks similar, but x86 element access mode is different

Data Structures in Assembly

Structures

- Alignment

Structs in C

- Way of defining compound data types
- A structured group of variables, possibly including other `structs`

```
struct song{
    int lengthInSeconds;
    int yearRecorded;
};


struct song song1;
song1.lengthInSeconds = 213;
song1.yearRecorded = 1994;

struct song song2;
song2.lengthInSeconds = 214;
song2.yearRecorded = 1988;
```

- Given a struct instance, access member using the `.` operator
- In assembly: pointer holds address of the first byte. Access elements with offsets.


Structure Representation

```
typedef struct rec {
    int a[4];
    long i;
    struct rec *next;
} *rec_p;
```



- **Characteristics:**
 - Contiguously-allocated region of memory
 - Refer to members within structure by names
 - Members may be of different types


```
typedef struct rec {
    int a[4];
    long i;
    struct rec *next;
} *rec_p;
```



- **Structure represented as block of memory:**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration order**
 - Even if another ordering would be more compact
- **Compiler determines overall size + positions of fields**
 - Machine-level programs has no understanding of the structures in the source code

Accessing a Structure Member

```
typedef struct rec {
    int a[4];
    long i;
    struct rec *next;
} *rec_p;
```




- Compiler knows the **offset** of each member within a struct
- Compute as
 - `* (r+offset)`
 - Referring to absolute offset, so no pointer arithmetic

```
long get_i(rec_p r)
{
    return r->i;
}
```

```
# r in %rdi
movq 16(%rdi), %rax
ret
```

Exercise: Pointer to Structure Member

```
typedef struct rec {
    int a[4];
    long i;
    struct rec *next;
} *rec_p;
```



```
long get_i(rec_p r){
    return r->i;
}
```


```
# r in %rdi
movq 16(%rdi), %rax
ret
```

```
rec_p get_next(rec_p r){
    return r->next;
}
```

```
# r in %rdi
movq 24(%rdi), %rax
ret
```

Generating Pointer to Array Element

```
typedef struct rec {
    int a[4];
    long i;
    struct rec *next;
} *rec_p;
```



■ Generating Pointer to Array Element:

- Offset of each structure member determined at compile time
- Compute as: $r + 4 * \text{index}$

```
int find_addr_of_array_elem
(rec_p r, long index){
    return r->a[index];
}
```

```
# r in %rdi, index in %rsi
movl    (%rdi,%rsi,4), %eax
ret
```

Memory Alignment in x86-64

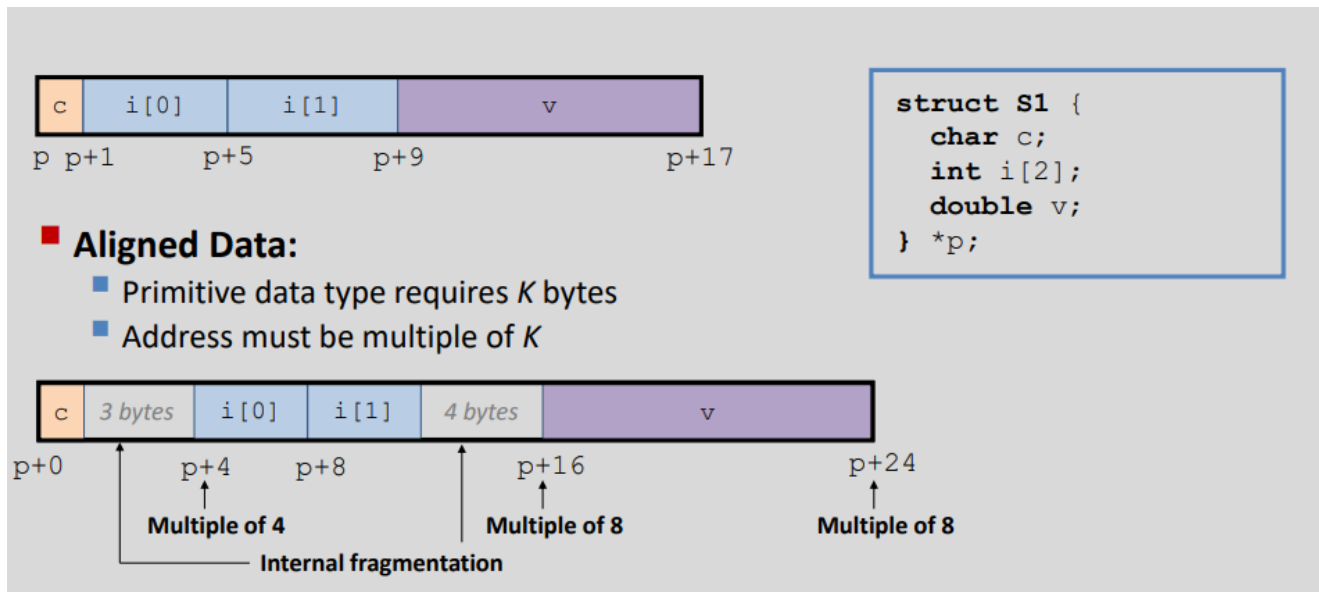
- For good memory system performance, Intel recommends data to be aligned
 - However, the x86-64 hardware will work correctly regardless of alignment of data
- *Aligned* means that any primitive object of K bytes must have an address that is multiple of K

K	Type	Addresses
1	char	No restrictions
2	short	Lowest bit must be zero: $\dots 0_2$
4	int, float	Lowest 2 bits zero: $\dots 00_2$
8	long, double	Lowest 3 bits zero: $\dots 000_2$
16	long double	Lowest 4 bits zero: $\dots 0000_2$

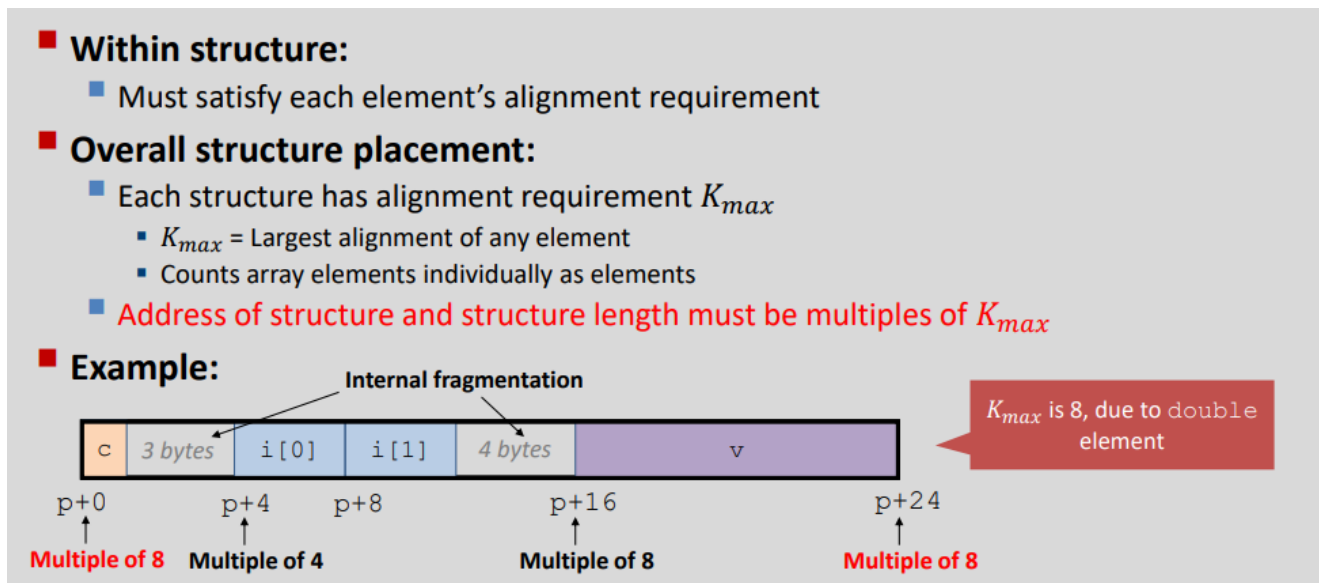
Alignment Principles

- **Aligned Data**
 - Primitive data type requires K bytes
 - Address must be multiple of K
 - Required on some machines; advised on x86-64
- **Motivation for Aligning Data**
 - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store value that spans word boundaries

Structures and Alignment



Satisfying Alignment with Structures




Alignment of Structs

■ Compiler will do the following:

- Maintains declared **ordering** of fields in struct
- Each **field** must be aligned **within** the struct (*may insert padding*)
 - `offsetof` can be used to get actual field offset
- Overall struct must be **aligned** according to largest field
- Total struct **size** must be multiple of its alignment (*may insert padding*)
 - `sizeof` should be used to get true size of structs

Accessing Array's Struct Elements

- Compute start of array element as: $12 * \text{index}$
 - `sizeof(S3) = 12`, including alignment padding
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`



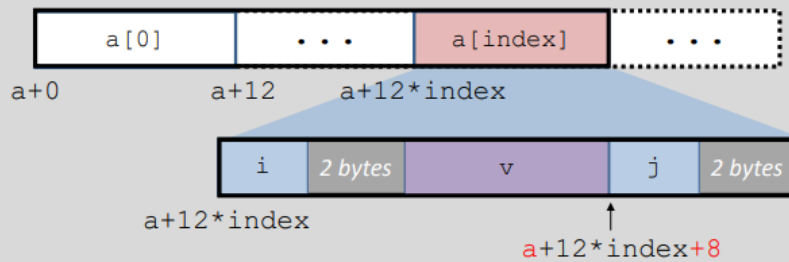
```
struct S3 {
    short i;
    int v;
    short j;
} a[10];
```

```
short get_j(int index){
    return a[index].j;
};
```

```
# %rdi = index
leaq (%rdi, %rdi, 2), %rax
movzwl a+8(,%rax,4), %eax
```

Accessing Array's Struct Elements

- Compute start of array element as: $12 * \text{index}$
 - `sizeof(S3) = 12`, including alignment padding
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`



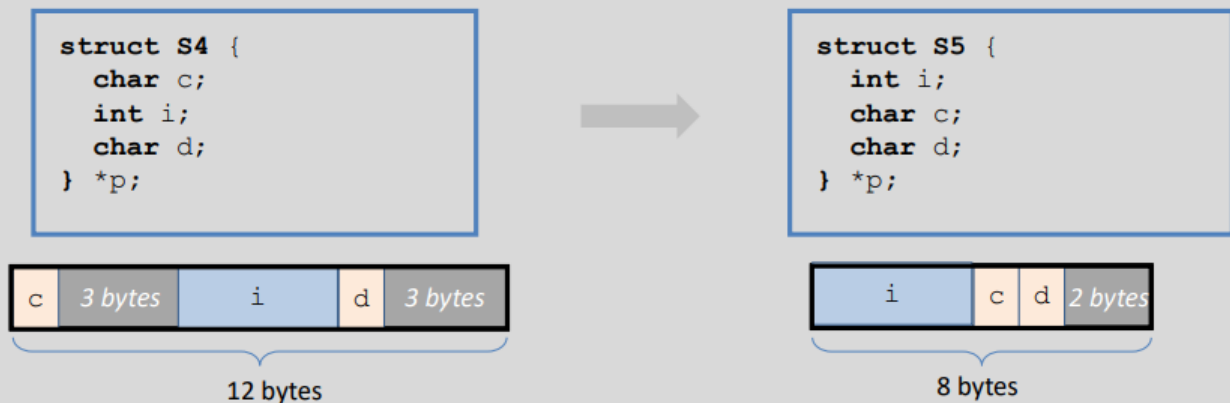
```
struct S3 {
    short i;
    int v;
    short j;
} a[10];
```

```
short get_j(int index){
    return a[index].j;
};
```

```
# %rdi = index
leaq (%rdi, %rdi, 2), %rax
movzwl a+8(,%rax,4), %eax
```

How the Programmer Can Save Space

- The compiler must respect the order elements are declared in
 - Sometimes the programmer can save space by declaring large data types first



Summary

■ Arrays

- Contiguous allocations of memory
- Can usually be treated like a pointer to first element
- Nested arrays
 - all levels in one contiguous block of memory
- Multi-Level arrays
 - First level in one contiguous block of memory
 - Each element in the first level points to another “sub” array
 - Parts anywhere in memory

■ Structures

- Allocate bytes in order declared
- Padding in middle and at end to satisfy alignment